

```
In [1]:  
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)  
import matplotlib.pyplot as plt # this is used for the plot the graph  
import numpy as np  
import seaborn as sns  
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
```

```
In [2]:  
avgPrice = pd.read_csv('data/BrandAverageRetailPrice.csv')  
avgPrice=avgPrice.rename(columns={'Brands':'Brand'})  
avgPrice=avgPrice.rename(columns={"vs. Prior Period":"ARP_pp"})  
avgPrice.head()
```

```
Out[2]:  
      Brand    Months      ARP      ARP_pp  
0 #BlackSeries 08/2020  15.684913      NaN  
1 #BlackSeries 09/2020      NaN -1.000000  
2 #BlackSeries 01/2021  13.611428      NaN  
3 #BlackSeries 02/2021  11.873182 -0.127705  
4 #BlackSeries 03/2021      NaN -1.000000
```

```
In [3]:  
brandDetails=pd.read_csv('data/BrandDetails.csv')  
brandDetails.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 144977 entries, 0 to 144976  
Data columns (total 25 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   State            144977 non-null   object    
 1   Channel          144977 non-null   object    
 2   Category L1      144977 non-null   object    
 3   Category L2      144977 non-null   object    
 4   Category L3      144245 non-null   object    
 5   Category L4      102618 non-null   object    
 6   Category L5      50135 non-null   object    
 7   Brand             144977 non-null   object    
 8   Product Description 144977 non-null   object    
 9   Total Sales ($)   144977 non-null   object    
 10  Total Units       144977 non-null   object    
 11  ARP               144977 non-null   float64   
 12  Flavor             7807 non-null    object    
 13  Items Per Pack    144977 non-null   int64    
 14  Item Weight        64454 non-null   object    
 15  Total THC           144977 non-null   object    
 16  Total CBD           144977 non-null   object    
 17  Contains CBD        144977 non-null   object    
 18  Pax Filter          44301 non-null   object    
 19  Strain              115639 non-null   object    
 20  Is Flavored         11287 non-null   object    
 21  Mood Effect          144977 non-null   object    
 22  Generic Vendor       144977 non-null   object    
 23  Generic Items         144977 non-null   object    
 24  $5 Price Increment  144977 non-null   object    
dtypes: float64(1), int64(1), object(23)  
memory usage: 27.7+ MB
```

```
In [4]:  
brandTotalSales=pd.read_csv('data/BrandTotalSales.csv')  
brandTotalSales.head()
```

```
Out[4]:  
      Months     Brand  Total Sales ($)  
0  09/2018  10x Infused  1,711.334232  
1  09/2018  1964 Supply Co.  25,475.21594500000  
2  09/2018  3 Bros Grow  120,153.644757  
3  09/2018      3 Leaf  6,063.5297850000000  
4  09/2018      350 Fire  631,510.0481550000
```

```
In [5]:  
brandTotalUnits=pd.read_csv('data/BrandTotalUnits.csv')  
brandTotalUnits=brandTotalUnits.rename(columns={'Brands':'Brand'})  
brandTotalUnits=brandTotalUnits.rename(columns={"vs. Prior Period":"totalUnits_pp"})  
brandTotalUnits.head()
```

```
Out[5]:  
      Brand    Months      Total Units  totalUnits_pp  
0 #BlackSeries 08/2020  1,616.3390040000000      NaN  
1 #BlackSeries 09/2020      NaN -1.000000  
2 #BlackSeries 01/2021  715.5328380000000      NaN  
3 #BlackSeries 02/2021  766.669135  0.071466  
4 #BlackSeries 03/2021      NaN -1.000000
```

```
In [6]:  
prodSales=pd.read_csv('data/Top50ProductsbyTotalSales-Timeseries.csv')  
prodSales.head()
```

Out[6]:

	Products	Months	Total Sales (\$)
0	Flower - Strain Blends - Flower (Gram)	07/2020	22,738,489.622206017
1	Flower - Strain Blends - Flower (Gram)	03/2021	22,648,507.64839804
2	Flower - Strain Blends - Flower (Gram)	05/2021	22,338,755.88508607
3	Flower - Strain Blends - Flower (Gram)	09/2020	21,461,950.605336975
4	Flower - Strain Blends - Flower (Gram)	06/2021	21,347,569.064233065

In [7]:

```
#convert months to datetime, convert total sales to int
prodSales['Months'] = pd.to_datetime(prodSales['Months'])
prodSales['Total Sales ($)'] = prodSales['Total Sales ($)').str.replace(',', '').astype(float)
```

In [8]:

```
#pull out only features of interest from brand details
newBrandDetails[['Product Description', 'Brand', 'Category L1', 'Category L2', 'Category L3', 'Category L4', 'Category L5',
                 'Flavor', 'Item Weight', 'Total THC', 'Total CBD', 'Pax Filter', 'Strain', 'Mood Effect',
                 'Generic Vendor', 'Generic Items']]
```

newBD=new.copy()

In [9]:

```
#convert total THC, total CBD to int
newBD['Total THC'] = newBD['Total THC'].str.replace(',', '').astype(float)
newBD['Total CBD'] = newBD['Total CBD'].str.replace(',', '').astype(float)
```

In [10]:

```
#convert Pax Filter, Mood Effect, Generic Vendor, and Generic Items to binary
```

```
def convertToBinaryPax(s):
    if s=='Pax':
        s=1
        s=float(s)
    if s=='Not Pax':
        s=0
        s=float(s)
    return(s)

def convertToBinaryMood(s):
    if type(s)==str and s=='Mood Specific':
        s=1
        s=float(s)
    if type(s)==str and s=='Not Mood Specific':
        s=0
        s=float(s)
    return(s)

def convertToBinaryVendor(s):
    if type(s)==str and s=='Generic Vendors':
        s=1
        s=float(s)
    if type(s)==str and s=='Non-Generic Vendors':
        s=0
        s=float(s)
    return(s)

def convertToBinaryItems(s):
    if type(s)==str and s=='Generic Items':
        s=1
        s=float(s)
    if type(s)==str and s=='Non-Generic Items':
        s=0
        s=float(s)
    return(s)

def convertToIntWeight(s):
    if type(s)!=float and 'mg' in s:
        s=s.replace('mg','')
        s=float(s)
        s=s/1000
    elif type(s)!=float:
        s=float(s)
    return(s)

newBD['Pax Filter'] = newBD['Pax Filter'].apply(convertToBinaryPax)
newBD['Mood Effect'] = newBD['Mood Effect'].apply(convertToBinaryMood)
newBD['Generic Vendor'] = newBD['Generic Vendor'].apply(convertToBinaryVendor)
newBD['Generic Items'] = newBD['Generic Items'].apply(convertToBinaryItems)
newBD['Item Weight'] = newBD['Item Weight'].apply(convertToIntWeight)
```

In [11]:

prodSales.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1671 entries, 0 to 1670
Data columns (total 3 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Products         1671 non-null   object 
 1   Months           1671 non-null   datetime64[ns]
 2   Total Sales ($)  1671 non-null   float64
dtypes: datetime64[ns](1), float64(1), object(1)
memory usage: 39.3+ KB
```

In [12]:

newBD.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144977 entries, 0 to 144976
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Product Description  144977 non-null   object 
 1   Brand              144977 non-null   object 
 2   Category L1         144977 non-null   object 
 3   Category L2         144977 non-null   object 
 4   Category L3         144245 non-null   object 
 5   Category L4         102618 non-null   object 
 6   Category L5         50135 non-null    object  
 7   Flavor              7807 non-null    object  
 8   Item Weight         64454 non-null   float64
 9   Total THC            144977 non-null   float64
 10  Total CBD             144977 non-null   float64
 11  Pax Filter           44301 non-null   float64
 12  Strain              115639 non-null   object  
 13  Mood Effect          144977 non-null   float64
 14  Generic Vendor       144977 non-null   float64
 15  Generic Items         144977 non-null   float64
dtypes: float64(7), object(9)
memory usage: 17.7+ MB

```

```
In [13]: #merge newBD and prodSales by product
newBD= newBD.rename(columns={'Product Description': 'Products'})
```

```
In [14]: merge= pd.merge(prodSales, newBD, on='Products', how='left')
```

```
In [15]: #product is not a key for the brand description values and some products are present more than once because they
# have different pax filter/category L3/category L5 values. This doesn't make any sense in our merged dataset
# because we don't know which product the total sales refers too. We will drop the repeated products from our
# dataset
new_merge = merge.drop_duplicates(subset=['Products','Months'], ignore_index=True)
```

```
In [16]: new_merge.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1671 entries, 0 to 1670
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Products          1671 non-null   object 
 1   Months             1671 non-null   datetime64[ns] 
 2   Total Sales ($)    1671 non-null   float64
 3   Brand              1671 non-null   object  
 4   Category L1        1671 non-null   object 
 5   Category L2        1671 non-null   object 
 6   Category L3        1671 non-null   object 
 7   Category L4        955 non-null    object  
 8   Category L5        360 non-null    object  
 9   Flavor              451 non-null    object  
 10  Item Weight         360 non-null   float64
 11  Total THC            1671 non-null   float64
 12  Total CBD             1671 non-null   float64
 13  Pax Filter           216 non-null   float64
 14  Strain              572 non-null   object  
 15  Mood Effect          1671 non-null   float64
 16  Generic Vendor       1671 non-null   float64
 17  Generic Items         1671 non-null   float64
dtypes: datetime64[ns](1), float64(8), object(9)
memory usage: 235.1+ KB

```

```
In [17]: #convert date column into separate month/year input
new_merge.loc[:, 'year']=new_merge.loc[:, 'Months'].dt.year
new_merge.loc[:, 'month']=new_merge.loc[:, 'Months'].dt.month
```

```
/Users/noelawheeler/opt/anaconda3/lib/python3.9/site-packages/pandas/core/indexing.py:1667: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self.obj[key] = value
```

```
In [18]: time_merge=new_merge.sort_values(by=['Months'])
time_merge.head(100)
```

Dat	Products	Months	Total Sales (\$)	Brand	Category L1	Category L2	Category L3	Category L4	Category L5	Flavor	Item Weight	Total THC	Total CBD	Pax Filter	Strain	Mood Effect	Generic Vendor
1385	Up North Humboldt - Durban Poison - Flower (Gram)	2018-10-01	232056.738077	Up North Humboldt	Inhaleables	Flower	Sativa	NaN	NaN	NaN	NaN	0.0	0.0	NaN	Durban Poison	0.0	0.0
398	Plus Products - Gummies - Indica - Unwind - Bl..	2018-10-01	717699.472242	Plus Products	Ingestibles	Edibles	Candy	Gummie Candy	NaN	Blackberry Lemon	NaN	90.0	10.0	NaN	NaN	1.0	0.0
1503	Caliva - Alien OG - Flower	2018-10-01	172868.662357	Caliva	Inhaleables	Flower	Hybrid	NaN	NaN	NaN	NaN	0.0	0.0	NaN	Alien OG	0.0	0.0

In [19]:

```

df_time=new_merge.copy()
products=new_merge['Products'].unique()
for i in products:
    y=time_merge.index[df_time['Products']==i].tolist()
    df_time.loc[y,'Rolling Average'] = (df_time.loc[y,'Total Sales ($)').shift(1) + df_time.loc[y,'Total Sales ($)].shift(2) + df_time.loc[y,'Total Sales ($)]
    df_time.at[y[2],'Rolling Average'] = (df_time.iloc[y[0]]['Total Sales ($)'] + df_time.iloc[y[1]]['Total Sales ($)'])/2
    df_time.at[y[1],'Rolling Average'] = df_time.iloc[y[0]]['Total Sales ($)']
    df_time.at[y[0],'Rolling Average'] = df_time.iloc[y[0]]['Total Sales ($)']

```

In [20]:

df_time.head()

Out[20]:

	Products	Months	Total Sales (\$)	Brand	Category L1	Category L2	Category L3	Category L4	Category L5	Flavor	...	Total THC	Total CBD	Pax Filter	Strain	Mood Effect	Generic Vendor	Generic Items	year	month	R Aw
0	Flower - Strain Blends - Flower (Gram)	2020-07-01	2.273849e+07	Flower	Inhaleables	Flower	Hybrid	NaN	NaN	NaN	...	0.0	0.0	NaN	NaN	0.0	1.0	1.0	2020	7	4.406626
1	Flower - Strain Blends - Flower (Gram)	2021-03-01	2.264851e+07	Flower	Inhaleables	Flower	Hybrid	NaN	NaN	NaN	...	0.0	0.0	NaN	NaN	0.0	1.0	1.0	2021	3	2.918791
2	Flower - Strain Blends - Flower (Gram)	2021-05-01	2.233876e+07	Flower	Inhaleables	Flower	Hybrid	NaN	NaN	NaN	...	0.0	0.0	NaN	NaN	0.0	1.0	1.0	2021	5	2.468878
3	Flower - Strain Blends - Flower (Gram)	2020-09-01	2.146195e+07	Flower	Inhaleables	Flower	Hybrid	NaN	NaN	NaN	...	0.0	0.0	NaN	NaN	0.0	1.0	1.0	2020	9	7.910952
4	Flower - Strain Blends - Flower (Gram)	2021-06-01	2.134757e+07	Flower	Inhaleables	Flower	Hybrid	NaN	NaN	NaN	...	0.0	0.0	NaN	NaN	0.0	1.0	1.0	2021	6	6.606598

5 rows × 21 columns

In [21]:


df_time.hist(figsize=(20,15))




Out[21]:

```

array([{'

```






```
In [22]: #we can see that pax filter is not very useful because it only has one value in our data set
#we can also drop generic vendor because it appear to have the same distributions as generic items
df_dropped=df_time.drop(['Pax Filter'], axis=1)
```

```
In [50]: corr_matrix=df_dropped.corr()
corr_matrix['Total Sales ($)']
```

```
Dat[50]:
Total Sales ($)      1.000000
Item Weight       -0.189561
Total THC        -0.143469
Total CBD         -0.080867
Mood Effect      -0.082969
Generic Vendor   0.453812
Generic Items     0.453812
year              0.008976
month             0.006425
Rolling Average  0.046546
Name: Total Sales ($), dtype: float64
```

```
In [25]: heatmap=sns.heatmap(corr_matrix, annot=True)
```



```
In [27]: from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import FunctionTransformer
from sklearn.base import BaseEstimator, TransformerMixin
```

```
In [51]: #augmentation: count how many products are offered by a brand
#drop Month column, we will look at year/month individually
df_dropped['Prod Counts']=df_dropped.groupby(['Brand'])['Products'].transform('count')
df_dropped=df_dropped.drop(columns=['Months'])
df_dropped.head()
```

```
Dat[51]:
   Products  Total Sales ($)  Brand  Category L1  Category L2  Category L3  Category L4  Category L5  Flavor  Item Weight  Total THC  Total CBD  Strain  Mood Effect  Generic Vendor  Generic Items  year  month  Rolling Average  Count
0  Flower - Strain  2.273849e+07  Flower  Inhaleables  Flower  Hybrid  NaN  NaN  NaN  0.0  0.0  NaN  0.0  1.0  1.0  2020  7  4.406626e+05  1
  Blends - Flower (Gram)
1  Flower - Strain  2.264851e+07  Flower  Inhaleables  Flower  Hybrid  NaN  NaN  NaN  0.0  0.0  NaN  0.0  1.0  1.0  2021  3  2.918791e+05  1
  Blends - Flower (Gram)
2  Flower - Strain  2.233876e+07  Flower  Inhaleables  Flower  Hybrid  NaN  NaN  NaN  0.0  0.0  NaN  0.0  1.0  1.0  2021  5  2.468878e+05  1
  Blends - Flower (Gram)
```

```
In [29]: num_pipeline=Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler())
])

cat_pipeline=Pipeline([
    ('imputer', SimpleImputer(strategy="most_frequent")),
    ('OneHotEncoder', OneHotEncoder(categories='auto', handle_unknown='ignore'))
])

numerical=["Total THC", "Total CBD", "Item Weight", "year", "month"]
categorical=['Products', 'Brand', 'Category L1', 'Category L2', 'Category L3', 'Category L4', 'Category L5', 'Flavor', 'Strain', 'Mood Effect', 'Generic Items']

full_pipeline=ColumnTransformer([
    ('cat', cat_pipeline, categorical),
    ('num', num_pipeline, numerical)
])
```

```
In [30]: target_column=df_dropped['Total Sales ($)']
df=df_dropped.drop(columns=['Total Sales ($)'])
```

```
In [31]: x=full_pipeline.fit_transform(df)
x_df = pd.DataFrame(x.toarray())
```

```
In [32]: x_df.shape
```

```
Out[32]: (1671, 149)
```

```
In [33]: y=target_column
```

```
In [43]: x_train, x_test, y_train, y_test = train_test_split(x_df, y, test_size=0.15)
```

```
In [35]: import statsmodels.api as sm
# build the OLS model (ordinary least squares) from the training data
sales_stats = sm.OLS(target_column, x_df)

# do the fit and save regression info (parameters, etc) in results_stats
results_stats = sales_stats.fit()
print(results_stats.summary())
```

OLS Regression Results

	Dep. Variable:	Total Sales (\$)	R-squared:	0.981	
	Model:	OLS	Adj. R-squared:	0.981	
	Method:	Least Squares	F-statistic:	1656.	
	Date:	Sat, 04 Dec 2021	Prob (F-statistic):	0.00	
	Time:	16:58:26	Log-Likelihood:	-23862.	
No. Observations:	1671	AIC:	4.783e+04		
Df Residuals:	1619	BIC:	4.811e+04		
Df Model:	51				
Covariance Type:	nonrobust				
	coef	std err	t	P> t	
				[0.025 0.975]	
0	-7.024e+05	2.03e+04	-34.596	0.000	-7.42e+05 -6.63e+05
1	5.279e+04	3.79e+04	1.393	0.164	-2.15e+04 1.27e+05
2	-5.667e+04	3.6e+04	-1.573	0.116	-1.27e+05 1.4e+04
3	-7.049e+05	3.29e+04	-21.437	0.000	-7.69e+05 -6.4e+05
4	-2.778e+05	1.66e+04	-16.726	0.000	-3.1e+05 -2.45e+05
5	-5.175e+05	1.51e+04	-34.286	0.000	-5.47e+05 -4.88e+05
6	-4.062e+05	2.83e+04	-14.349	0.000	-4.62e+05 -3.51e+05
7	-4.105e+05	2.86e+04	-14.334	0.000	-4.67e+05 -3.54e+05
8	-4.553e+05	2.77e+04	-16.431	0.000	-5.1e+05 -4.01e+05
9	-1.579e+06	2.76e+04	-57.160	0.000	-1.63e+06 -1.52e+06
10	-2.246e+06	2.76e+04	-81.324	0.000	-2.3e+06 -2.19e+06
11	7.718e+06	3.58e+04	215.469	0.000	7.65e+06 7.79e+06
12	1.537e+05	1.89e+04	8.147	0.000	1.17e+05 1.91e+05
13	3.247e+05	2.51e+04	12.963	0.000	2.76e+05 3.74e+05
14	-3.124e+05	8138.983	-38.384	0.000	-3.28e+05 -2.96e+05
15	1.908e+05	3.84e+04	4.962	0.000	1.15e+05 2.66e+05
16	-6.409e+04	3.09e+04	-2.073	0.038	-1.25e+05 -3453.660
17	-1.183e+05	3.09e+04	-3.829	0.000	-1.79e+05 -5.77e+04
18	-5.55e+04	3.34e+04	-1.662	0.097	-1.21e+05 1e+04
19	-1.619e+04	3.09e+04	-0.524	0.601	-7.68e+04 4.45e+04
20	-1.099e+05	3.84e+04	-2.859	0.004	-1.85e+05 -3.45e+04
21	7.614e+04	2.35e+04	3.239	0.001	3e+04 1.22e+05
22	1.824e+04	2.35e+04	0.776	0.438	-2.79e+04 6.44e+04
23	1.314e+04	4.65e+04	0.283	0.777	-7.8e+04 1.04e+05
24	-5059.7078	4.65e+04	-0.109	0.913	-9.62e+04 8.61e+04
25	-5.717e+05	3.15e+04	-18.124	0.000	-6.34e+05 -5.1e+05
26	5.032e+05	2.22e+04	22.685	0.000	4.6e+05 5.47e+05
27	-5.012e+05	3.2e+04	-15.648	0.000	-5.64e+05 -4.38e+05
28	-1.577e+04	3.72e+04	-0.424	0.672	-8.87e+04 5.72e+04
29	-2.22e+04	3.87e+04	-0.573	0.566	-9.81e+04 5.37e+04
30	-2.4e+04	3.8e+04	-0.632	0.527	-9.85e+04 5.05e+04
31	3.03e+04	1.82e+04	1.662	0.097	-5458.396 6.61e+04
32	2.677e+04	2.89e+04	0.926	0.355	-3e+04 8.35e+04
33	-2.482e+04	2.6e+04	-0.955	0.339	-7.58e+04 2.61e+04
34	-1.085e+04	1.47e+04	-0.739	0.460	-3.96e+04 1.79e+04

35	-6.731e+05	2.4e+04	-28.020	0.000	-7.2e+05	-6.26e+05
36	-7.016e+05	2.03e+04	-34.558	0.000	-7.41e+05	-6.62e+05
37	-5.105e+05	4.1e+04	-12.467	0.000	-5.91e+05	-4.3e+05
38	5.792e+05	2.77e+04	20.944	0.000	5.25e+05	6.33e+05
39	5.922e+05	2.77e+04	21.418	0.000	5.38e+05	6.46e+05
40	4.852e+05	2.55e+04	19.038	0.000	4.35e+05	5.35e+05
41	2.04e+05	2.4e+04	8.509	0.000	1.57e+05	2.51e+05
42	2.101e+05	2.4e+04	8.763	0.000	1.63e+05	2.57e+05
43	-7.114e+05	2.17e+04	-32.802	0.000	-7.54e+05	-6.69e+05
44	3.047e+05	1.5e+04	20.303	0.000	2.75e+05	3.34e+05
45	-6.021e+04	3.34e+04	-1.801	0.072	-1.26e+05	5374.752
46	1.052e+05	5.06e+04	2.081	0.038	6024.133	2.04e+05
47	-8.374e+04	3.27e+04	-2.564	0.010	-1.48e+05	-1.97e+04
48	8.826e+04	3.27e+04	2.702	0.007	2.42e+04	1.52e+05
49	6.248e+04	3.27e+04	1.913	0.056	-1581.415	1.27e+05
50	-7.024e+05	2.03e+04	-34.596	0.000	-7.42e+05	-6.63e+05
51	-3878.8384	1.75e+04	-0.222	0.825	-3.82e+04	3.04e+04
52	-1.5e+06	2.25e+04	-66.550	0.000	-1.54e+06	-1.46e+06
53	-1.272e+06	2.17e+04	-58.506	0.000	-1.31e+06	-1.23e+06
54	3.893e+06	1.91e+04	204.072	0.000	3.86e+06	3.93e+06
55	1.537e+05	1.89e+04	8.147	0.000	1.17e+05	1.91e+05
56	3.247e+05	2.51e+04	12.963	0.000	2.76e+05	3.74e+05
57	-3.124e+05	8138.983	-38.384	0.000	-3.28e+05	-2.96e+05
58	-7.882e+04	1.73e+04	-4.558	0.000	-1.13e+05	-4.49e+04
59	8080.8052	1.2e+04	0.673	0.501	-1.55e+04	3.16e+04
60	-5.697e+05	2.38e+04	-23.938	0.000	-6.16e+05	-5.23e+05
61	-3.167e+04	1.75e+04	-1.805	0.071	-6.61e+04	2739.166
62	1951.8776	1.99e+04	0.098	0.922	-3.72e+04	4.11e+04
63	-1.085e+04	1.47e+04	-0.739	0.460	-3.96e+04	1.79e+04
64	-6.731e+05	2.4e+04	-28.020	0.000	-7.2e+05	-6.26e+05
65	-7.016e+05	2.03e+04	-34.558	0.000	-7.41e+05	-6.62e+05
66	1.146e+06	2.27e+04	50.411	0.000	1.1e+06	1.19e+06
67	4.14e+05	1.32e+04	31.457	0.000	3.88e+05	4.4e+05
68	-7.114e+05	2.17e+04	-32.802	0.000	-7.54e+05	-6.69e+05
69	3.047e+05	1.5e+04	20.303	0.000	2.75e+05	3.34e+05
70	1.12e+05	1.61e+04	6.961	0.000	8.05e+04	1.44e+05
71	-1.879e+05	8051.292	-23.332	0.000	-2.04e+05	-1.72e+05
72	6.965e+04	1.02e+04	6.859	0.000	4.97e+04	8.96e+04
73	-2.922e+04	1.32e+04	-2.219	0.027	-5.5e+04	-3389.752
74	-6.197e+04	9830.213	-6.304	0.000	-8.12e+04	-4.27e+04
75	-5.146e+05	1.07e+04	-47.991	0.000	-5.36e+05	-4.94e+05
76	-1.879e+05	8051.292	-23.332	0.000	-2.04e+05	-1.72e+05
77	4.323e+04	1.07e+04	4.042	0.000	2.23e+04	6.42e+04
78	3.946e+05	1.19e+04	33.112	0.000	3.71e+05	4.18e+05
79	9.078e+04	1.26e+04	7.207	0.000	6.61e+04	1.15e+05
80	2.642e+04	1.04e+04	2.534	0.011	5968.202	4.69e+04
81	-6.197e+04	9830.213	-6.304	0.000	-8.12e+04	-4.27e+04
82	-6.197e+04	9830.213	-6.304	0.000	-8.12e+04	-4.27e+04
83	-5.923e+04	1.03e+04	-5.750	0.000	-7.94e+04	-3.9e+04
84	9.438e+04	9243.165	10.211	0.000	7.63e+04	1.13e+05
85	-2.778e+05	1.66e+04	-16.726	0.000	-3.1e+05	-2.45e+05
86	3.258e+06	2.17e+04	149.839	0.000	3.22e+06	3.3e+06
87	6.569e+05	1.96e+04	33.593	0.000	6.19e+05	6.95e+05
88	-1.579e+06	2.76e+04	-57.160	0.000	-1.63e+06	-1.52e+06
89	1.016e+05	1.25e+04	8.159	0.000	7.72e+04	1.26e+05
90	8080.8052	1.2e+04	0.673	0.501	-1.55e+04	3.16e+04
91	-1.085e+04	1.47e+04	-0.739	0.460	-3.96e+04	1.79e+04
92	3.047e+05	1.5e+04	20.303	0.000	2.75e+05	3.34e+05
93	-2.246e+06	2.76e+04	-81.324	0.000	-2.3e+06	-2.19e+06
94	2.642e+04	1.04e+04	2.534	0.011	5968.202	4.69e+04
95	-2.369e+05	1.32e+04	-17.999	0.000	-2.63e+05	-2.11e+05
96	-1.879e+05	8051.292	-23.332	0.000	-2.04e+05	-1.72e+05
97	9.438e+04	9243.165	10.211	0.000	7.63e+04	1.13e+05
98	2.642e+04	1.04e+04	2.534	0.011	5968.202	4.69e+04
99	3.642e+05	1.03e+04	35.516	0.000	3.44e+05	3.84e+05
100	-2.778e+05	1.66e+04	-16.726	0.000	-3.1e+05	-2.45e+05
101	-1.879e+05	8051.292	-23.332	0.000	-2.04e+05	-1.72e+05
102	8080.8052	1.2e+04	0.673	0.501	-1.55e+04	3.16e+04
103	2.807e+05	1.64e+04	17.141	0.000	2.49e+05	3.13e+05
104	-5.175e+05	1.51e+04	-34.286	0.000	-5.47e+05	-4.88e+05
105	9.438e+04	9243.165	10.211	0.000	7.63e+04	1.13e+05
106	4.016e+05	1.31e+04	30.662	0.000	3.76e+05	4.27e+05
107	-5.175e+05	1.51e+04	-34.286	0.000	-5.47e+05	-4.88e+05
108	-1.879e+05	8051.292	-23.332	0.000	-2.04e+05	-1.72e+05
109	-2.178e+05	1.94e+04	-11.216	0.000	-2.56e+05	-1.8e+05
110	7.614e+04	2.35e+04	3.239	0.001	3e+04	1.22e+05
111	1.052e+05	5.06e+04	2.081	0.038	6024.133	2.04e+05
112	1.824e+04	2.35e+04	0.776	0.438	-2.79e+04	6.44e+04
113	-8.374e+04	3.27e+04	-2.564	0.010	-1.48e+05	-1.97e+04
114	8.826e+04	3.27e+04	2.702	0.007	2.42e+04	1.52e+05
115	1.908e+05	3.84e+04	4.962	0.000	1.15e+05	2.66e+05
116	-6.409e+04	3.09e+04	-2.073	0.038	-1.25e+05	-3453.660
117	-6.021e+04	3.34e+04	-1.801	0.072	-1.26e+05	5374.752
118	6.248e+04	3.27e+04	1.913	0.056	-1581.415	1.27e+05
119	-1.183e+05	3.09e+04	-3.829	0.000	-1.79e+05	-5.77e+04
120	-2.482e+04	2.6e+04	-0.955	0.339	-7.58e+04	2.61e+04
121	-5.55e+04	3.34e+04	-1.662	0.097	-1.21e+05	1e+04
122	-1.619e+04	3.09e+04	-0.524	0.601	-7.68e+04	4.45e+04
123	-1.099e+05	3.84e+04	-2.859	0.004	-1.85e+05	-3.45e+04
124	-5.717e+05	3.15e+04	-18.124	0.000	-6.34e+05	-5.1e+05
125	-7.024e+05	2.03e+04	-34.596	0.000	-7.42e+05	-6.63e+05
126	-4.062e+05	2.83e+04	-14.349	0.000	-4.62e+05	-3.51e+05
127	1.704e+06	1.9e+04	89.536	0.000	1.67e+06	1.74e+06
128	-7.114e+05	2.17e+04	-32.802	0.000	-7.54e+05	-6.69e+05
129	3.047e+05	1.5e+04	20.303	0.000	2.75e+05	3.34e+05
130	-4.105e+05	2.86e+04	-14.334	0.000	-4.67e+05	-3.54e+05
131	-4.553e+05	2.77e+04	-16.431	0.000	-5.1e+05	-4.01e+05
132	2.04e+05	2.4e+04	8.509	0.000	1.57e+05	2.51e+05
133	5.792e+05	2.77e+04	20.944	0.000	5.25e+05	6.33e+05
134	5.032e+05	2.22e+04	22.685	0.000	4.6e+05	5.47e+05
135	-7.016e+05	2.03e+04	-34.558	0.000	-7.41e+05	-6.62e+05
136	1.537e+05	1.89e+04	8.147	0.000	1.17e+05	1.91e+05
137	5.922e+05	2.77e+04	21.418	0.000	5.38e+05	6.46e+05
138	2.101e+05	2.4e+04	8.763	0.000	1.63e+05	2.57e+05
139	-5.012e+05	3.2e+04	-15.648	0.000	-5.64e+05	-4.38e+05

```

140      -6.843e+04   1.24e+04    -5.516     0.000   -9.28e+04   -4.41e+04
141      -1.409e+05   1.27e+04   -11.086     0.000   -1.66e+05   -1.16e+05
142      -1.606e+06   1.13e+04   -142.186    0.000   -1.63e+06   -1.58e+06
143      1.397e+06   1.38e+04   101.207    0.000   1.37e+06   1.42e+06
144      4455.2043   1.6e+04     0.278     0.781   -2.69e+04   3.58e+04
145      2.997e+04   1.44e+04     2.088     0.037   1814.764   5.81e+04
146      -4.563e+05   9323.573   -48.939    0.000   -4.75e+05   -4.38e+05
147      6.981e+04   1.04e+04     6.734     0.000   4.95e+04   9.01e+04
148      3.392e+04   1.02e+04     3.328     0.001   1.39e+04   5.39e+04
=====
Omnibus:                      585.560 Durbin-Watson:           0.584
Prob(Omnibus):                 0.000 Jarque-Bera (JB):       75278.771
Skew:                           0.559 Prob(JB):                  0.00
Kurtosis:                      35.863 Cond. No.            8.97e+16
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 9.63e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [58]: from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
```

```
In [36]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(x_train, y_train)
predictions = lr.predict(x_test)
```

```
In [37]: from sklearn.metrics import mean_absolute_error

preds = lr.predict(x_test)
rmse = mean_absolute_error(y_test, preds)
rmse
```

```
Out[37]: 232942.94321512
```

```
In [283]: x_df.shape
```

```
Out[283]: (1671, 149)
```

```
In [38]: #PCA Analysis
from sklearn import decomposition
pca = decomposition.PCA(n_components=4)

pca_df = pca.fit_transform(x_df)
```

```
In [39]: pca_df.shape
```

```
Out[39]: (1671, 4)
```

```
In [40]: #new testing/training data using pca
x_train, x_test, y_train, y_test = train_test_split(pca_df, y, test_size=0.15)
```

```
In [44]: #use random forest ensemble method
from sklearn.ensemble import RandomForestRegressor

rf=RandomForestRegressor(n_estimators=20, random_state=0)
rf.fit(x_train, y_train.ravel())
pred=rf.predict(x_test)
```

```
In [38]: #PCA Analysis
from sklearn import decomposition
pca = decomposition.PCA(n_components=4)

pca_df = pca.fit_transform(x_df)
```

```
In [39]: pca_df.shape
```

```
Out[39]: (1671, 4)
```

```
In [40]: #new testing/training data using pca
x_train, x_test, y_train, y_test = train_test_split(pca_df, y, test_size=0.15)
```

```
In [44]: #use random forest ensemble method
from sklearn.ensemble import RandomForestRegressor

rf=RandomForestRegressor(n_estimators=20, random_state=0)
rf.fit(x_train, y_train.ravel())
pred=rf.predict(x_test)
```

```
In [45]: rmse = mean_absolute_error(y_test, preds)
rmse
```

```
Out[45]: 1317184.6205840847
```

```
In [46]: from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
```

```
In [62]: #Employ Kfold cross validation to training results
from sklearn.model_selection import KFold
from sklearn import model_selection

#KFold, with 10 splits where we first shuffle our data before splitting it
kfolds = model_selection.KFold(n_splits=10, random_state=42, shuffle=True)
rf_model_kfold = RandomForestRegressor(n_estimators=20, random_state=0)

lr_model_kfold = LinearRegression()

rf_results_kfold = model_selection.cross_val_score(rf_model_kfold, x_df, y.ravel(), scoring='r2', cv=kfolds)
lr_results_kfold = model_selection.cross_val_score(lr_model_kfold, x_df, y.ravel(), scoring='r2', cv=kfolds)

print("Random Forest Accuracy: %.2f%%" % (rf_results_kfold.mean()*100.0))
print("Linear Regression Accuracy: %.2f%%" % (lr_results_kfold.mean()*100.0))

Random Forest Accuracy: 99.08%
Linear Regression Accuracy: 79.85%
```

```
In [35]: #gridsearch method to optimize parameters
from sklearn.model_selection import GridSearchCV
parameters = {'criterion':('squared_error', 'absolute_error', 'poisson'), 'max_depth':[5, 10, 15, 20]}
randforest=RandomForestRegressor()
clf = GridSearchCV(randforest, parameters)
clf.fit(x_df, target_column)
```

```
Out[35]: GridSearchCV(estimator=RandomForestRegressor(),
param_grid={'criterion': ('squared_error', 'absolute_error',
'poisson'),
'max_depth': [5, 10, 15, 20]})
```

```
In [314]: print(clf.best_params_)

{'criterion': 'poisson', 'max_depth': 15}
```

```
In [48]: #experiment with other models
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn import svm

gpr = GaussianProcessRegressor(random_state=0)
gpr.fit(x_train, y_train.ravel())
pred_gpr=gpr.predict(x_test)

svm_model=svm.SVR()
svm_model.fit(x_train, y_train.ravel())
pred_svm=svm_model.predict(x_test)
```

```
In [49]: rmse_gpr = mean_absolute_error(y_test, pred_gpr)
rmse_svm = mean_absolute_error(y_test, pred_svm)
print("RMSE Gaussian Process Regressor: " + str(rmse_gpr))
print("RMSE SVM Regressor: " + str(rmse_svm))

RMSE Gaussian Process Regressor: 1720693.592204542
RMSE SVM Regressor: 742022.4251756461
```

```
In [ ]:
```